

STA 141A - Spring 2025 - Homework 4

Instructor: Dr. Akira Horiguchi

Student name: ABCDE FGHIJ; Student ID: 123456789

Due date: May 7, 2025 at 9 PM (PT)

The assignment must be done in an [R Markdown](#) or [Quarto](#) document. The assignment must be submitted by the due date above by uploading:

1. a .pdf file in GRADESCOPE (if you can knit/compile your .rmd to a .html file only, please save the created .html file as a .pdf file (by opening the .html file -> print -> save to .pdf)).

Email submissions will not be accepted.

Each answer has to be based on R code that shows how the result was obtained. The code has to answer the question or solve the task. For example, if you are asked to find the largest entry of a vector, the code has to return the largest element of the vector. If the code just prints all values of the vector, and you determine the largest element by hand, this will not be accepted as an answer. No points will be given for answers that are not based on R. This homework already contains chunks for your solution (you can also create additional chunks for each solution if needed, but it must be clear to which tasks your chunks belong).

There are many possible ways to write R code that is needed to answer the questions or do the tasks, but for some of the questions or tasks you might have to use something that has not been discussed during the lectures or the discussion sessions. You will have to come up with a solution on your own. Try to understand what you need to do to complete the task or to answer the question, feel free to search the Internet for possible solutions, and discuss possible solutions with other students. It is perfectly fine to ask what kind of an approach or a function other students use. However, you are not allowed to share your code or your answers with other students. Everyone has to write the code, do the tasks and answer the questions on their own.

During the discussion sessions, you may be asked to present and share your solutions.

Also, for your information, you can write python code chunks in RMarkdown or Quarto documents. For example, the following code chunk shows how to print the numbers from 0 to 3 in Python:

```
for i in range(4):  
    print("Current iteration:", i)
```

```
## Current iteration: 0  
## Current iteration: 1  
## Current iteration: 2  
## Current iteration: 3
```

(You will need Python already installed to run the following code chunk. If this causes an error and you don't plan to use Python, feel free to delete this chunk and stick with R for this course.)

1. Linear Regression with a generated dataset

```
GenerateData_v1 <- function(n) {  
  x <- runif(n, min=-1, max=1)  
  y <- -2 + 3 * x + rnorm(n, sd = 0.5)  
  data.frame(x = x, y = y)  
}  
test_data <- GenerateData_v1(n=100000)
```

(a) Generate a training dataset with $n = 10$ points using the function `GenerateData_v1()`. Fit a simple linear regression model to this dataset. Compute the training MSE. Compute the test MSE using the dataset `test_data`. How do the two MSEs compare? *Hint*: you may want to use the functions `lm()` and `predict()`.

Your Solution (Code)

Your Solution (Text)

(b) Do the same as in part (a), but now by generating a training dataset with $n = 100$ points.

Your Solution (Code)

Your Solution (Text)

(c) Do the same as in part (a), but now by generating a training dataset with $n = 1000$ points.

Your Solution (Code)

Your Solution (Text)

(d) As n increases, what value do the test MSEs seem to converge to? How does this value relate to the noise standard deviation in the data generation process?

Your Solution (Text)

2. Linear Regression with the `anscombe` dataset

Consider the linear regression model

$$y = \beta_0 + \beta_1 X + \varepsilon, \quad \varepsilon \sim N(0, \sigma^2).$$

and the `anscombe` dataset, whose columns have names `x1`, `x2`, `x3`, `x4`, `y1`, `y2`, `y3`, `y4`.

With this dataset, fit a linear regression model for each of the four x - y pairs. For each pair, do/answer the following:

1. What are the estimated coefficients $\hat{\beta}_0$ and $\hat{\beta}_1$?
2. How about the R^2 ?
3. Plot y against x , along with the fitted regression line. For the plots, comment on your findings.

```
### Your Solution (Code)
```

Your Solution (Text)

3. Linear Regression with the mpg data set

(a) Load the mpg dataset from `library(ggplot2)`. Add new variables as follows, and show the frequency table of those new variables.

- A binary variable of the car types: value `M` if the car type is either 2seater, compact, or midsize, and value `L` if the car type is either minivan, pickup, subcompact, or suv
- A binary variable of the type of transmission: value `auto` for auto cars, and value `manual` for manual cars

Your Solution (Code)

(b) Fit a linear regression model using `cty` as the response. For the predictors, use the variables defined in (a), as well as `displ`, `year`, and `cyl`. Comment on your findings.

Your Solution (Code)

Your Solution (Text)

(c) Plot the residuals against the fitted values. What can you conclude from this?

Your Solution (Code)

Your Solution (Text)